

# PD FEM CODE 3

## GENERAL INFO

---

### COMPILER ISSUES:

---

#### INTEL COMPILER

Compilation with Intel compiler and -O3 flag generates a binary *femswe* which is segfaulting for cubed sphere with 5 levels. => FIXED

The problem was the allocation of arrays with a given number of array entries on the subroutine's stack:

```
REAL*8 :: psi(nwhateverx)
```

This allocates the data directly on the thread's stack. Since this stack system is by far more limited compared to the system-wide memory, this segfaults for larger problem sizes since the thread is running out of stack size.

Solution: Dynamically allocate and deallocate large arrays, since this stores these arrays on the HEAP instead of the stack.

TODO: preallocate temporary arrays used during simulation time steps to avoid malloc/free overheads in subroutine *allocateall*

#### GNU COMPILER:

Helps discovering out of boundary access or boundary mismatch detection with compile option *-fcheck=all* . Some lines were fixed with the help of this compiler option.

## NEW FILES:

---

**adjacency\_matrices.f90:** Write out files with points of adjacency matrices and cell traversal

Usage: *./adjacency\_matrices [cube/hex] [cells]*

The data can be plot with 'gnuplot' , generating png images, e.g.

*\$> gnuplot adjacencies\_hex\_0000010242.gnuplot*

**sfc\_optimize.f90:** Reorder cells along SFC curve

Usage: *./sfc\_optimize [cube/hex] [cells]*

Additional sfc index tables are appended to the gridmap files to lookup the “old” index.

## F77 TO F90

---

Converted gengrid\_hex.f to gengrid\_hex.f90

Bitreproducible with original code if using compile option -O0

## ADDED PROGRAM ARGUMENTS:

---

*./gengrid\_hex [levels]*

*./gengrid\_cube [levels]*

*./buildfem\_op [hex/cube] [cells]*

*./femswe [parameter file]*

## MULTIGRID ISSUES WITH SFC OPTIMIZATION

---

The multi grid restriction operator assumes a particular cell indexing scheme (in *buildop\_fem.f90*). After an SFC reordering, this scheme does not exist anymore. Therefore the *sfc\_optimize.f90* code writes additional lookup tables to the grid-data files. These tables are then used to lookup the non-SFC-ordered position and the SFC-ordered position.

*fNewFaceId*: lookup table with old cell id stored at table entry.

*fNewFaceIdInverse*: lookup table with new cell id at table entry.

E.g. determining the cell for a MG prolongation operation on a hexagonal cell is then given by

*fNewFaceIdInverse(fNewFaceId(if0, igrd), igrdp)*, see *buildop\_fem.f90*

## BIT-WISE REPRODUCIBILITY

---

- The cell-ordering of the output files *run\_\** had to be reconstructed based on *fNewFaceIdInverse*
- Bit-wise reproducible results generated for *hex\_10242* and *cube\_13824* without SFC optimization.

With SFC optimization not yet reproducible due to some operations in the MG.

### IMPLEMENTATION

---

- The SFC optimization program *sfc\_optimize.f90* expects a grid-file as a data input and writes the SFC reordered cell-data to the same file or to an additional file if specified as a second program parameter.
- The reordering is based on computing the 1D SFC index and reordering all points. There's space for optimization regarding this ordering by using another SFC (Hilbert, Peano, Sierpinski, etc.):

The function *computeSFCIndex* can be easily extended with other SFC curves.

- SFC 1D coordinate is computed with the Z-Curve [ SUBROUTINE *computeSFCIndex(X, o\_scalar)* in *sfc\_optimize.f90* ] based on the 3D cell's mid-point.
- See results: optimization for cells not enough and sometimes leads to less performance since the originally already existing locality between cells and edges/vertices is destroyed.

### CELLS/EDGES/NODES

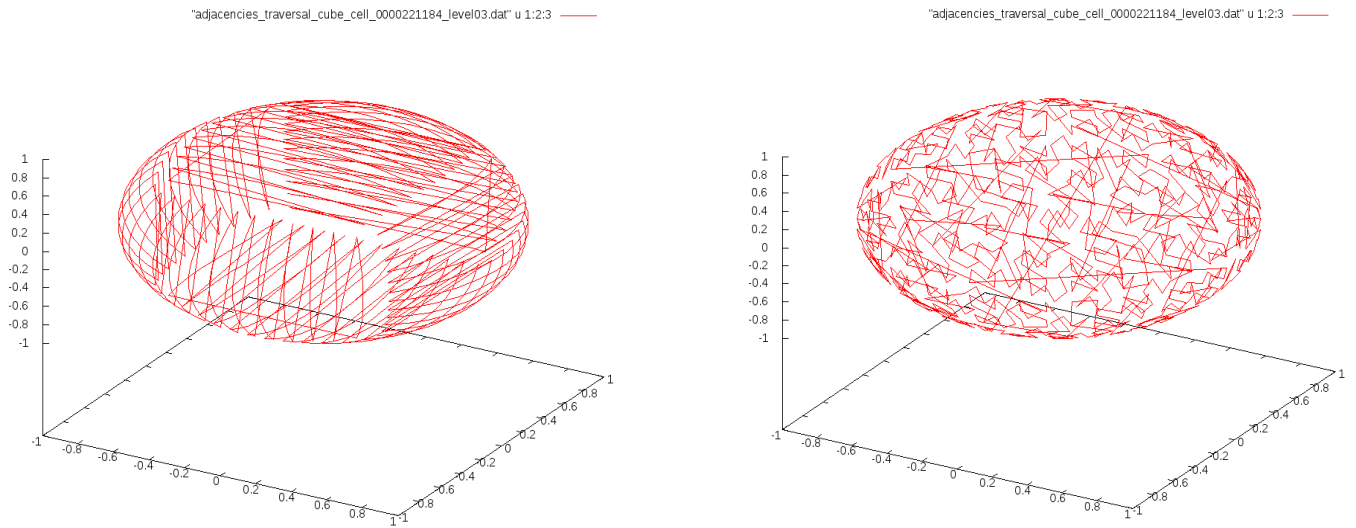
---

The SFC reordering was so far **only applied to the cells**.

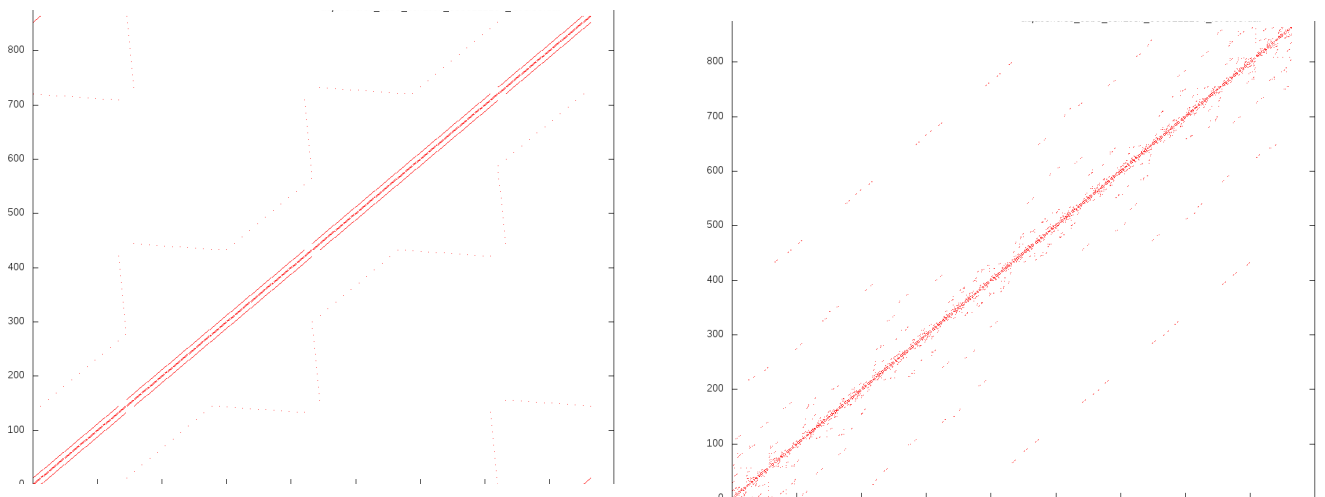
Also applying it to edges and nodes, we expect performance enhancements.

## CUBED SPHERE:

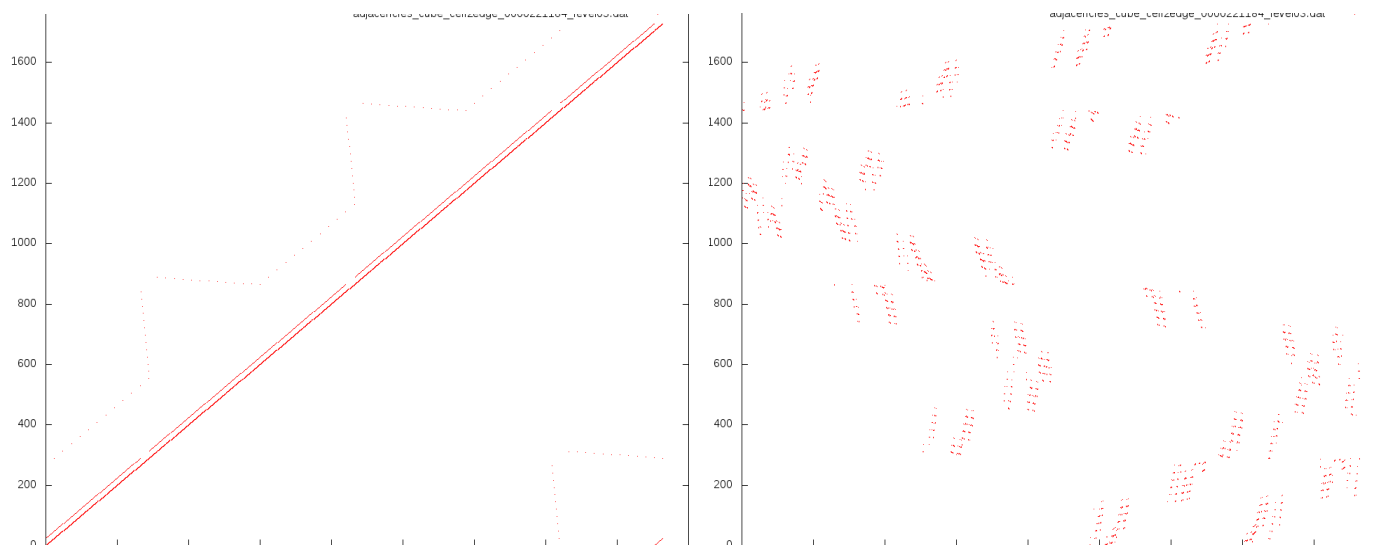
Displaying grid level 3 only



Left image: default ordering, right image: Z-ordered.



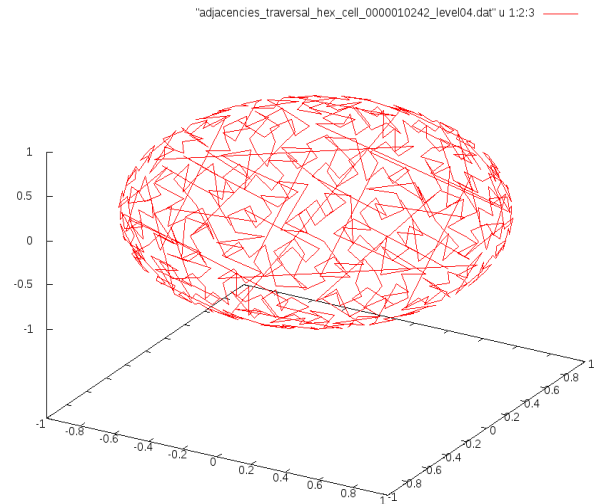
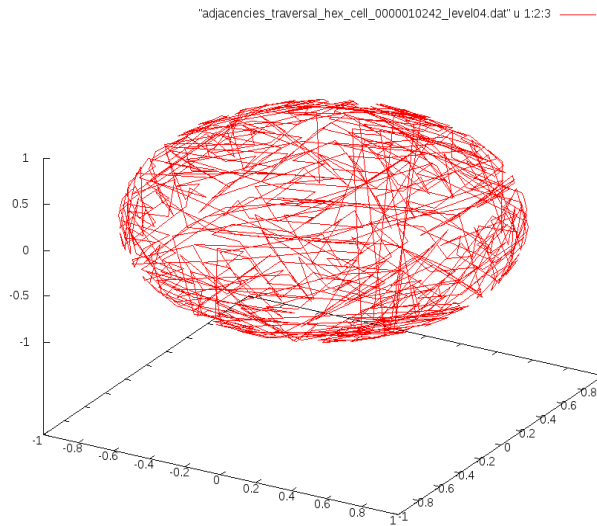
Left image: cell2cell adjacency matrix already almost diagonal. Right image: after ordering with the SFC.



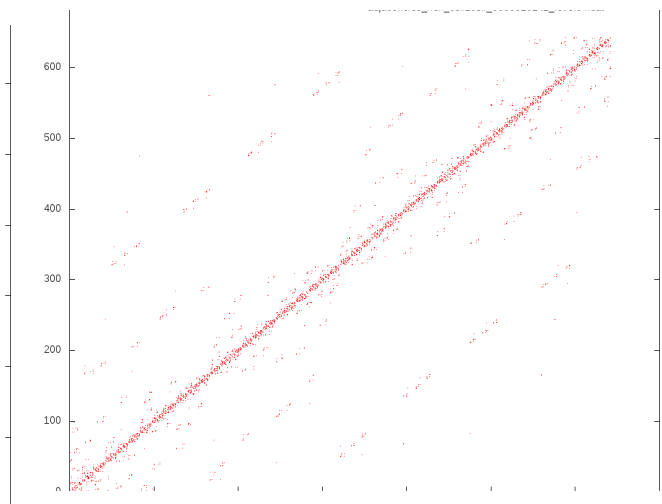
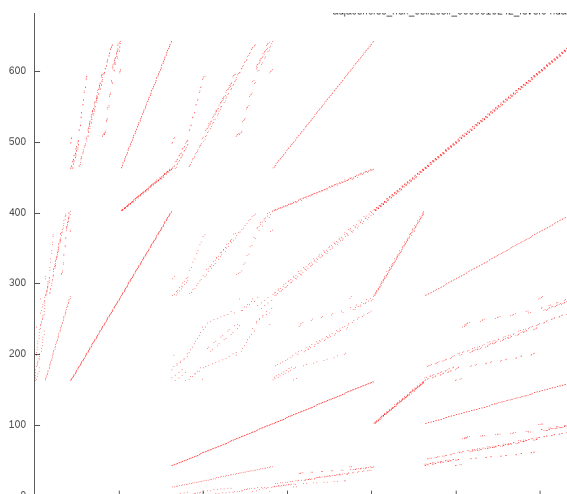
Reordering of cells “destroys” the edge-locality. Left image: Close-to diagonal cell-to-edge adjacency matrix. Right image: adjacency matrix after SFC-reordering

## HEX-GRIDS:

Info: Displaying grid-levels 4



















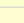
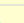




Left: default ordering of hex-grid. Right: Z-SFC ordered.



Left: Unordered cell2cell adjacency matrix, Right: SFC ordered

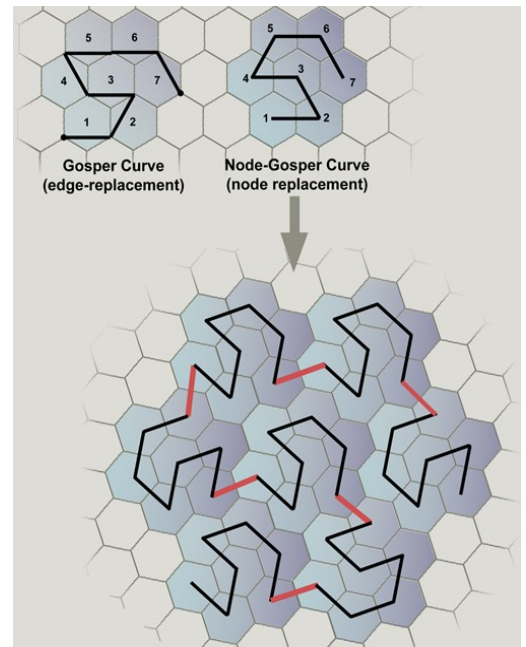
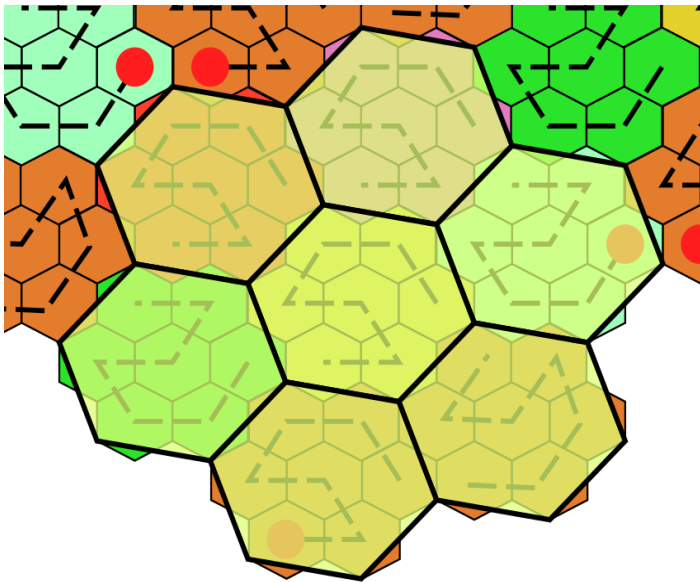
## Vtune Amplifier:

Function / Call Stack	CPU Time▼	★	Module	Function (Full)
±dualadvflx	195.863s		femswe	dualadvflx
±hodgeh	151.335s		femswe	hodgeh
±primaladvflx	126.612s		femswe	primaladvflx
±hodgej	28.428s		femswe	hodgej
±massm	22.053s		femswe	massm
±opert	12.660s		femswe	opert
±dprimal2	9.195s		femswe	dprimal2
±hodgehinv	8.846s		femswe	hodgehinv
±operw	8.818s		femswe	operw
±step	6.123s		femswe	step
±approxminv	5.165s		femswe	approxminv
±[import thunk __powidf2]	4.784s		femswe	[import thunk __powidf2]
±buildhelm	3.552s		femswe	buildhelm
±hodgejinv	3.100s		femswe	hodgejinv
±ddual1	2.323s		femswe	ddual1
±helmholtz	1.989s		femswe	helmholtz
±mgsolve	1.523s		femswe	mgsolve
±massl	1.503s		femswe	massl
±relax	1.463s		femswe	relax
±massminv	1.236s		femswe	massminv
±restrict	1.169s		femswe	restrict
±diagnostics	1.117s		femswe	diagnostics



- SFC ordering for **edges** and **vertices** missing
- Optimize array access such as `fnxtf(i,1:6,ngrids)` to `fnxtf(1:6,i,ngrids)` since data accessed via index 1:6 is consecutively aligned in memory -> optimizations for cacheline.
- Implement & test other SFC such as Gosper curve for hexagonal grids.

(from <http://courses.washington.edu/css342/fukuda/prog/prog3.html>)



- Adaptive Gosper SFCs:  
Hanging nodes & non-matching edges & other ugly properties

## PARALLELIZATION AND BENCHMARKS

---

### SOURCE CODE FILES:

Parallelization done with OpenMP for femswe.f90 only.

Other codes have to be parallelized as well. Otherwise, no large-scale simulations can be executed.

### MEMORY CONSUMPTION:

Info: Approx. 16 GB memory consumption for problem size with 655362 cells.

The reason for this excessive memory consumption can be e.g. found by allocating the same number of DOFs on each MG level. TODO: Further investigation needed.

### DEBUGGING:

Invalid OpenMP code detected with OpenMP debug version: Shared variables had to be private and were not optimized with OpenMP version => Always test with -O0, OpenMP and Intel.

### PUSH/PULL ACCESS SCHEME

=> Replaced push with pull calls for restriction operator

Bit reproducability tested for Cubed sphere with 13824 cells.

## USING WORKSHARE AND PARALLEL DO PARALLELIZATION:

- `!$OMP PARALLEL DO DEFAULT(NONE) PRIVATE(iv0) SHARED(nvertx, psi, igrd)`

`DO iv0 = 0, nvertx`

`psi(iv0) = .....`

`END DO`

`!$OMP END PARALLEL DO`

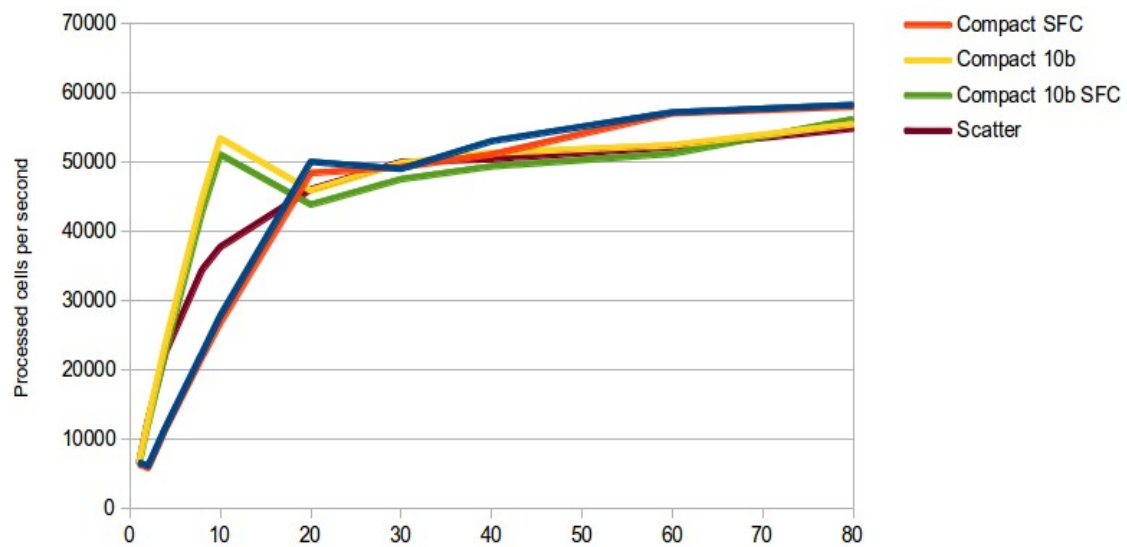
- Always use `DEFAULT(NONE)` to assure specification of `PRIVATE` or `SHARED` for all variables.

`PRIVATE`: The storage for these variables is replicated for each thread, thus independent to other threads. E.g. local counters or array indices inside the loop have to be declared as `PRIVATE`.

`SHARED`: The variables are directly accessed by the threads. Arrays should be typically declared as `SHARED` to avoid replicating them.

- Output of restart files and diagnostic computations deactivated.

- Benchmark:
  - Intel compiler environment variables
    - Scatter: export KMP\_AFFINITY=scatter
    - Compact: export KMP\_AFFINITY=compact (DO NOT USE THIS)
    - Compact 10: export KMP\_AFFINITY=fine,compact,1,0
  - Westmere 4 Socket, 10 cores per Socket, hyperthreaded: 80 threads available.



- Scalability up to 10 cores.
- No further scalability beyond 10 cores. Possible reasons:
  - First-touch policy on memory pages (close to bandwidth limited)?
  - Bandwidth limited?

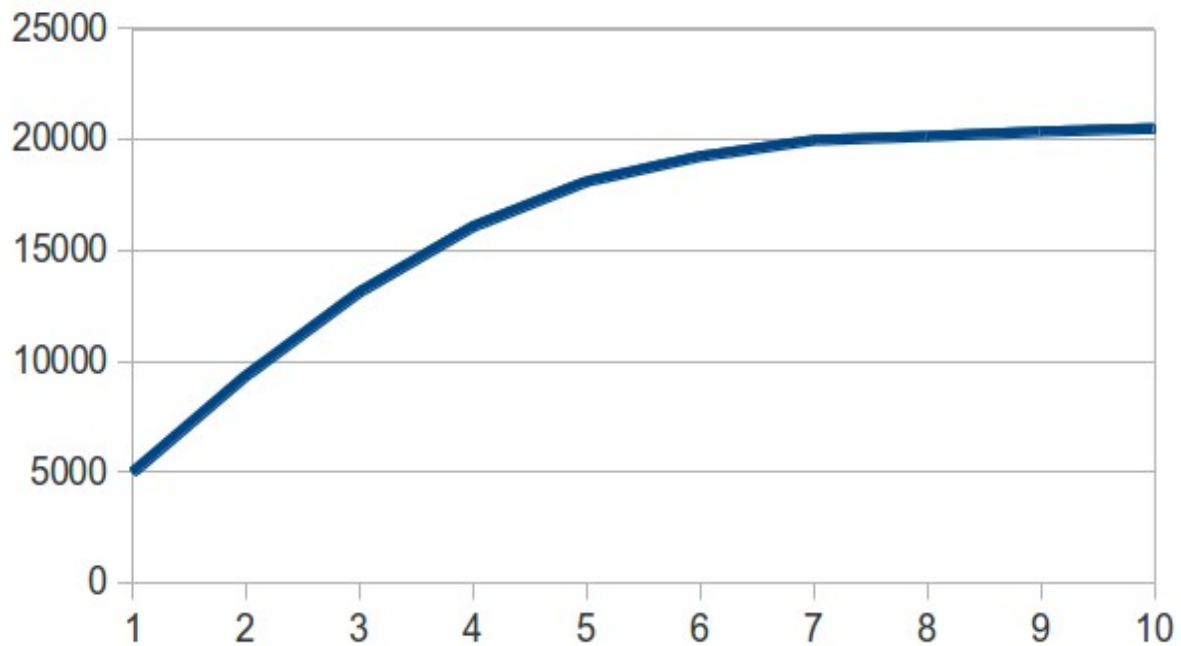
OR

  - Problem size is too small, overhead of parallelization too high?

## OPTIMIZING WITH FIRST-TOUCH POLICY FOR POSSIBLY MEMORY-BANDWIDTH LIMITED PROBLEM

Example: streaming benchmark which copies data from one memory location to another.

Throughput of streaming benchmark for 1 to 10 cores



Scalability limited for bandwidth limited problems starting at ~6 cores.

For multiple CPUs, there are also multiple memory controller.

If allocating the memory with a single core (e.g. by initializing the data with the first core only), all data gets allocated at the memory assigned to a single memory controller.

First touch policy can help to overcome these issues:

Initialize the data using all cores in parallel to allocate memory at memory controllers with the assumption that the main memory access is also accomplished via the same memory controller for computing a simulation timestep.

